



RIKEN's
Programs for
Junior Scientists

Automatic Retuning of Floating-Point Precision

Ivan R. Ivanov (TokyoTech) William S. Moses (UIUC)

ivanov.i.aa@m.titech.ac.jp, wsmoses@illinois.edu

Choosing the right floating point precision is important



RIKEN's
Programs for
Junior Scientists

Too high a precision

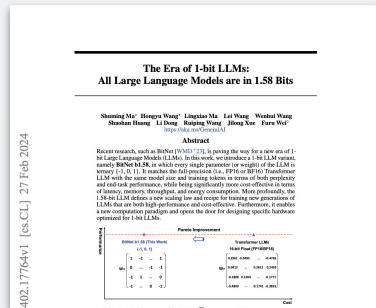





For nothing - if we didn't need the precision

- More power
- Longer runtime

Too low?

- Useless results
 - some scientific applications require > 128 bit fp numbers to work



The Herbie Project

Try • Install • Learn

Find and fix floating-point problems:

$$\text{sqrt}(x+1) - \text{sqrt}(x) \rightarrow 1/(\text{sqrt}(x+1) + \text{sqrt}(x))$$

Future hardware...?

Existing Approaches to Re-tuning or exploration



RIKEN's
Programs for
Junior Scientists

Macros (#define float float8)

- + Easy to set up
- Programs will often fail to build (wrong size memory to copy, cannot cast to new type, conflicting names)
- Applies to entire program, not just the sub-computation we care about

Templating functions

- Requires rewriting code to be generic (and breaks library calls)
- + Easier compiler-level errors, and more brute force debugging than endless conflicts
- + Fast
- No analysis

Adapt

- + Analysis of errors
- Requires rewriting code
- Slow
- Does not reflect fp optimized state

Valgrind?

Existing Approaches to Re-tuning



RIKEN's
Programs for
Junior Scientists

Macros (#define float float8)


- + Easy to set up
- Programs will often fail to build (wrong size memory to copy, cannot cast to new type, conflicting names)
- Applies to entire program, not just the sub-computation we care about

Templating functions

- Requires rewriting code to be generic (and breaks library calls)
- + Easier compiler-level errors, and more brute force debugging than endless conflicts
- + Fast
- No analysis

Adapt

- + Analysis of errors
- Requires rewriting code
- Slow
- **Does not reflect fp optimized state**



Valgrind?

FP Optimizations



RIKEN's
Programs for
Junior Scientists

```
define fun():  
    ...  
    %c = fadd %a, %b : f32  
    %e = fmul %d, %c : f32  
    ...
```



```
define fun():  
    ...  
    %e = fma %a, %b, %d : f32  
    ...
```

```
define fun():  
    ...  
    %c = call __fp_add(%a, %b) : %struct.fp  
    %e = call __fp_mul(%d, %c) : %struct.fp  
    ...
```



```
define fun():  
    ...  
    %c = call __fp_add(%a, %b) : %struct.fp  
    %e = call __fp_mul(%d, %c) : %struct.fp  
    ...
```

llvm: What is going on here???? I have no idea.

But what about optimization?



RIKEN's
Programs for
Junior Scientists

- Optimizations can change result!
Sometimes good, sometimes bad
 - Program can now run faster at expense of error (e.g. fastmath)
 - Optimizers (e.g. Herbie) can optimize for more accurate!



- Source-level tooling can fail to predict actual results.
- Even better tools would be able to jointly optimize the precision and the computation!

`-fassociative-math`

Allow re-association of operands in series of floating-point operations.

This allows the compiler to change the order of evaluation in a sequence of floating point operations. For example if you have an expression $(a + b) + c$, it can evaluate it instead as $a + (b + c)$. While these are mathematically equivalent with real numbers, they aren't equivalent in floating point arithmetic: the errors they incur can be different, in some cases quite significantly so:

```
julia> a = 1e9+1; b = -1e9; c = 0.1;

julia> (a+b)+c
1.1

julia> a+(b+c)
1.100000023841858
```

From Simon Byrne (<https://simonbyrne.github.io/notes/fastmath/>)

The Herbie Project



Try • Install • Learn

Find and fix floating-point problems:

$\text{sqrt}(x+1) - \text{sqrt}(x) \rightarrow 1/(\text{sqrt}(x+1) + \text{sqrt}(x))$

Impact of optimizations?



RIKEN's
Programs for
Junior Scientists

On FPBench: **745 functions** with fp computation

```
clang -O3 -ffast-math ...
```

vs

```
clang -O0
```


Mismatched results on **124** of them

Impact of optimizations?



RIKEN's
Programs for
Junior Scientists

Mismatched results on **124**/745 functions




You may get unexpected results on real hardware

Our Approach



RIKEN's
Programs for
Junior Scientists

- We build off the Enzyme infrastructure for generating shadow computations from generic LLVM code (historically derivatives) to now generate new functions which replace floating representations.
- Apply optimization prior (and optionally after) floating tuning!
- No program rewriting required, just call the corresponding intrinsic function. (Or specify full-module effect on command line)



Example Usage



RIKEN's
Programs for
Junior Scientists

```
#define FROM 64
#define TO_E 5    // Exponent
#define TO_M 21   // Mantissa

double f(double x, double y) {
    return ...;
}

double a = 3, b = 4;

// Step 1: Call wrapper to convert inputs into specified precision
//          Here we convert from fp64 to fp32, using the old type
//          as storage for the truncated data.
a = __enzyme_truncate_mem_value(a, FROM, TO_E, TO_M);
b = __enzyme_truncate_mem_value(b, FROM, TO_E, TO_M);

// Step 2: Generate an fp32 version of the fp64 function
//          with our new intrinsic.
double res = __enzyme_truncate_mem_func(f, FROM, TO_E, TO_M)(a, b);

// Step 3: Convert the truncated result back into the full fp64.
res = __enzyme_expand_mem_value(res, FROM, TO_E, TO_M);
```

Example Usage



RIKEN's
Programs for
Junior Scientists

```
#define FROM 64
#define TO_E 5    // Exponent
#define TO_M 21   // Mantissa

double f(double x, double y) {
    return ...;
}
```

```
double a = 3, b = 4;
```

```
// Step 1: Call w
//          Here w
//          as sto
a = __enzyme_trun
b = __enzyme_trun
```

```
// Step 2: Genera
//          with o
double res = __en
```

```
// Step 3: Conver
res = __enzyme_expand_mem_value(res, FROM, TO_E, TO_M);
```

You can also specify on the command line:

```
clang ... --enzyme-truncate-all="64to5-21"
```

For full-module effect

Future Work & Conclusions



RIKEN's
Programs for
Junior Scientists

- Being aware of the floating point representation used is critical to both performance and accuracy!
- Optimizing floating point choices inherently requires understanding the impact of normal compiler optimization
- Our tool builds off Enzyme/LLVM to automatically retune floats alongside optimization without programmer rewriting
- Future work:
 - Combine with program analysis to determine which operations requiring tuning, and by what amount
 - Tune only individual parts of a function but leave other parts at full precision
 - Consider other sources of error than floating point
- All open source (<https://github.com/EnzymeAD/Enzyme>), please reach out if interested!