

# Sign Extension Optimizations inside LLVM

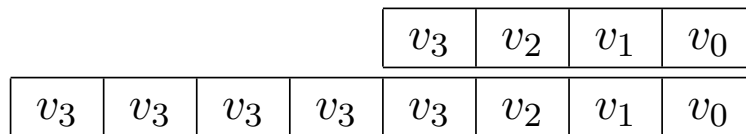
Panagiotis Karouzakis, Polyvios Pratikakis

University Of Crete, ICS-FORTH

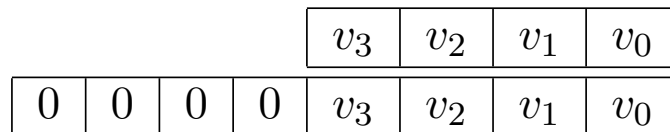
11/4/2024

# Extension Operations

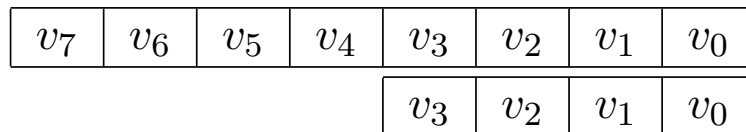
Sign extend the low 4 bits into a 8 bit value.



Zero extend the low 4 bits into a 8 bit value.



Truncate the 8 bit value into a 4 bit value. The high bits are lost.



# The Problem

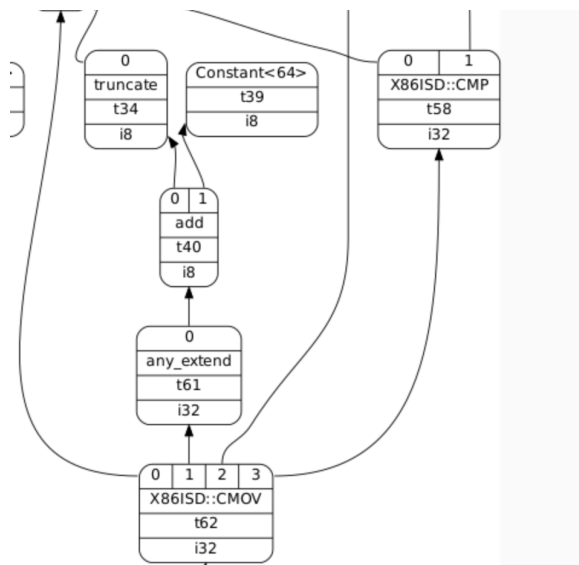


Figure 1: X86 DAG code after the last DAG combine

We could drop the `truncate` and the `any_extend` and do a `i32` addition. The `truncate` does not generate any code in X86 but the `any_extend` to `i32` does.

# The Problem

- In C/C++ code the type *int* may create many redundant sign extensions.
- Sign extension optimizations may need further research especially for RISC-V
- On X86 redundant truncations followed by `any_extend` are observed on SPEC 2017 benchmarks in many blocks.
- LLVM still doesn't query for all the legal widths that a operator has for a given target. It does that on a subset of the available operators.

- Kevin Redwine and Norman Ramsey[CC 2004] proposed the *filltypes*.
- They created type rules that used the *filltypes* to find the optimal solution.
- They used Dynamic Programming to find a solution.
- Their implementation was on the small Language C— a subset of C.
- They visited each AST node to apply all the type rules.

# Filltypes

- We use the notion of *filltype* proposed by Kevin Redwine et al[CC 2004 ]
- A *filltype* indicates what an operand produces and accepts in their upper bits.
  - Upper bits are the rest of the bits that do not have any data.
- An operand has a fill type only if it is *widenable*, i.e., if applying the operator to wide values can simulate the operator applied to narrow values.
- For example, since xor has a *filltype*,  $\text{xor}_{i32}$  can be implemented as  $\text{xor}_{i64}$  regardless of the high bits of the operands.

$$(\underline{1010}) \oplus (\underline{0011}) = 1001$$

$$(1111\underline{1010}) \oplus (1111\underline{0011}) = 0000\underline{1001}$$

4 bit xor implemented using 8 bit xor

# Filltypes

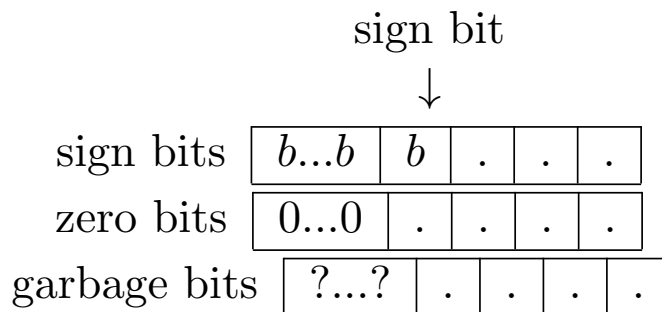
- We use symbols  $s$ ,  $z$ ,  $g$  to match sign, zero, and garbage upper bits respectively.
- The input operands must be typed before visiting an instruction and they can be an instruction as well.

$and :: g \times g \rightarrow g$

$and :: z \times g \rightarrow z$

$and :: g \times z \rightarrow z$

$and :: s \times s \rightarrow s$



- Create multiple solutions per Instruction that are legal for the target with different Instruction width.
- Use data flow information to learn how many bits are data.
- If we have upper bits left and not all of them are data we have a *filltype*.
- If not we can insert an extension to add a *filltype* based on the target
- If we insert a truncation we can potentially remove a *filltype*.



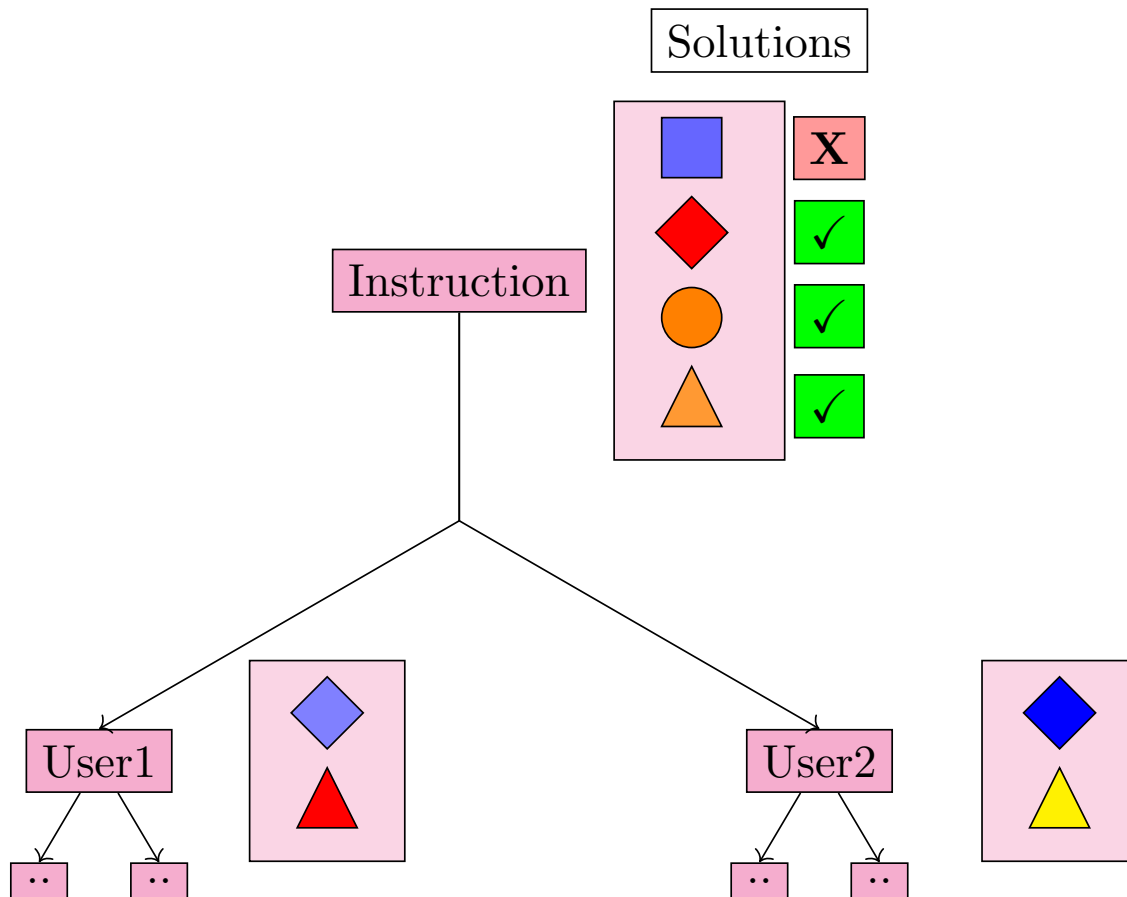
# Binary operators

- For binary operators search for legal *filltype* rules.
  - $xor :: s \times s \rightarrow s$
  - $xor :: z \times z \rightarrow z$
  - $xor :: g \times g \rightarrow g$
  - ...
- Ask Target Lowering for the legal Instructions widths.
- For example, many targets offer xor with 32 bits and 64 bits.
  - $xor :: 32 \times 32 \rightarrow 32$
  - $xor :: 64 \times 64 \rightarrow 64$
- If we have found an operation with legal Instruction width that has a *fillType*, we can create a solution that keeps the new width, the data bits and other information.

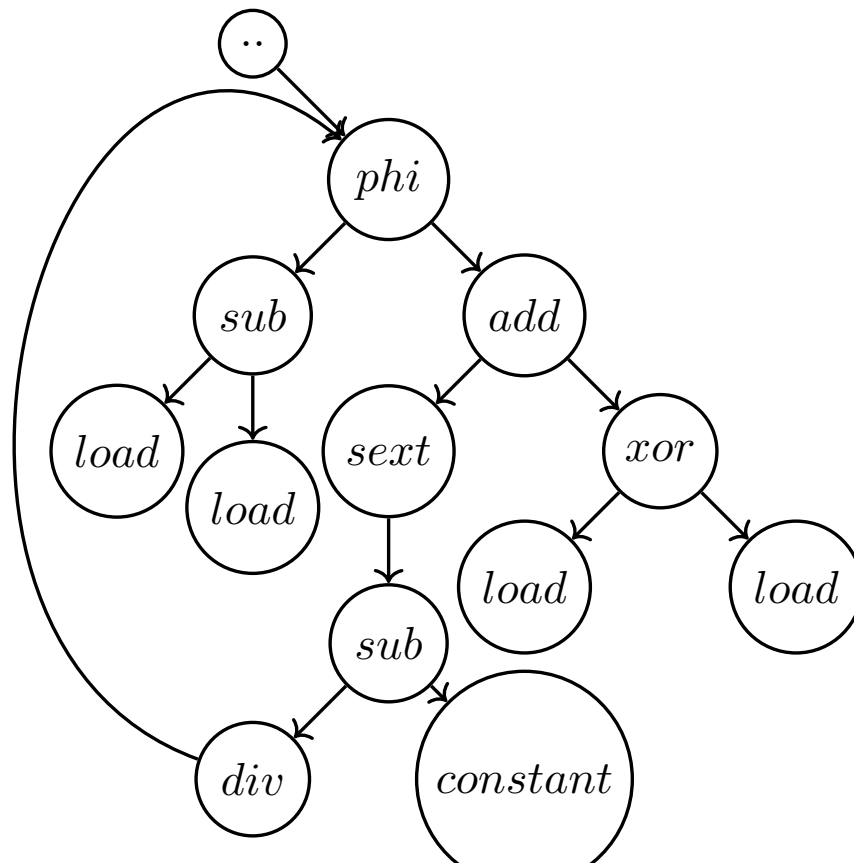
# Our approach

- Doing this optimization inside LLVM IR to add support for every Language Frontend.
- As a consequence we solve a flow sensitive problem to deal with the control flow, instead of a flow insensitive problem.
- We extend the proposed operand *filltypes* to match the LLVM operands.
- We have to deal with LLVM Intrinsics i.e., to choose among Intrinsics of different widths.
- While we use the available target operations, X86 needs special handling because some extensions and truncations are free.
- When we have more than 1 Use of an Instruction we might get conflicting solutions.

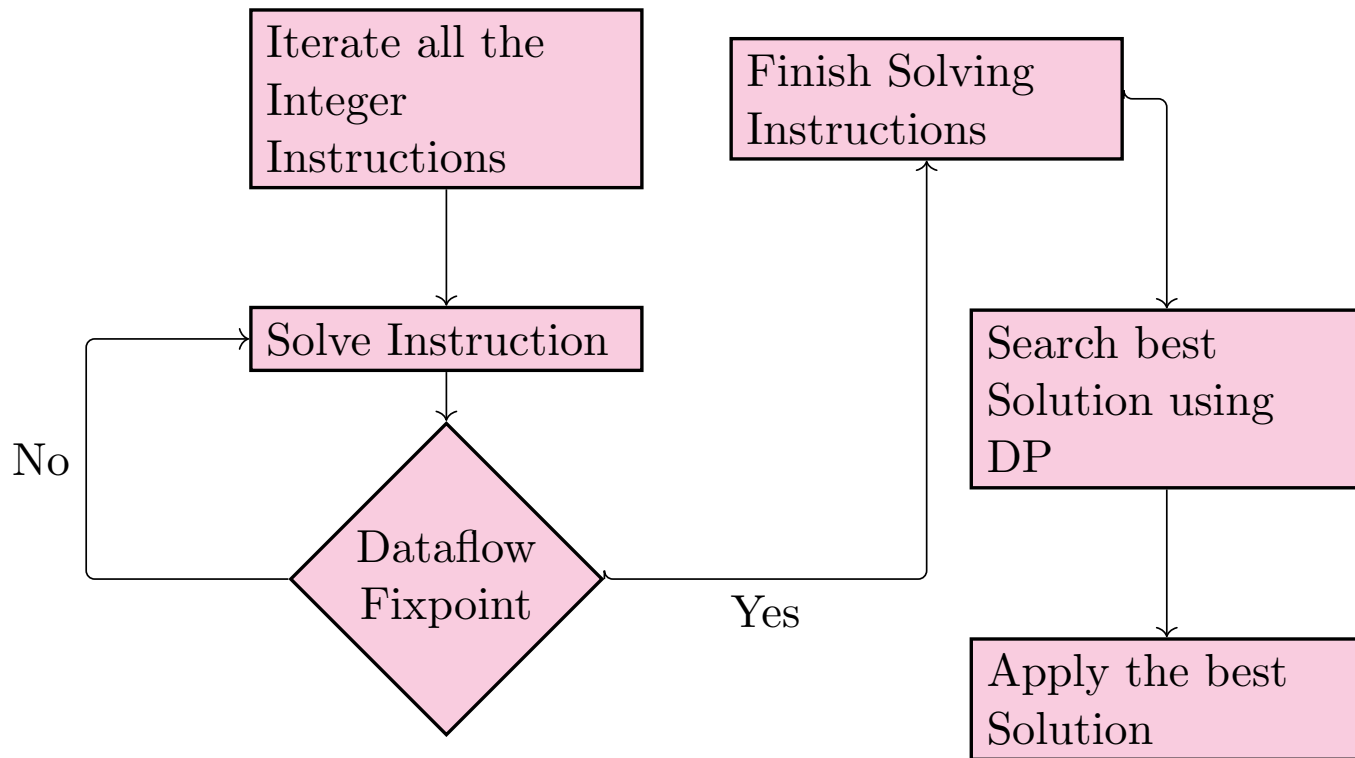
# Multiple Users Example



# PHIs example



# Implementation details



# Limitations and Next Steps

- If an operator overflows it requires special consideration.
- Checking for overflows is not one hundred percent accurate, so we lose optimization opportunities.
- Currently the project is implemented as a Function Pass. It may be useful to use Module pass to better infer function parameters.
- Using preferred X86 register width.
- Choosing between Intrinsics of different widths.